

Face-Mask-App: AI

November 25, 2020



Outline

- 1 Introduction
- 2 Model & Training
- 3 Datasets
 - Kaggle
 - Mask model
 - Face model
- 4 Model results
- 5 Model to mobile

Introduction

- Neural networks
- Two different models
 - Face model
 - Mask model
- Transfer learning
 - ResNet18
- PyTorch

Model & Training

Model

- Pretty standard transfer learning model
 - Log Softmax output
 - Negative Log Likelihood loss function
 - Stochastic Gradient Descent with Momentum
 - Cyclical learning rate
- Pretty standard stuff!

```
def get_model(dataloaders, n_epochs=30):  
    device = 'cuda' if torch.cuda.is_available() else 'cpu'  
  
    # load pretrained resnet model  
    model = models.resnet18(pretrained=True)  
  
    # replace output layer, we have two outputs  
    model.fc = nn.Sequential(  
        nn.Linear(model.fc.in_features, 2),  
        nn.LogSoftmax(dim=1)  
    )  
  
    # transform to GPU if needed  
    model = model.to(device)  
  
    # loss function  
    criterion = nn.NLLLoss()  
  
    # Learning rate and momentum will be overridden by scheduler  
    optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)  
  
    # One Cycle Policy scheduler  
    scheduler = torch.optim.lr_scheduler.OneCycleLR(  
        optimizer,
```

Model & Training

Training

- Standard PyTorch training loop
 - Resize images
 - Random data augmentation
- Nothing special here!

```
# Iterate over data.
for inputs, labels in dataloaders[phase]:
    inputs = inputs.to(device)
    labels = labels.to(device)

    # zero the parameter gradients
    optimizer.zero_grad()

    # forward
    # track history if only in train
    with torch.set_grad_enabled(phase == 'train'):
        outputs = model(inputs)
        loss = criterion(outputs, labels)

    # backward + optimize only if in training phase
    if phase == 'train':
        loss.backward()
        optimizer.step()
        scheduler.step()

# statistics
_, preds = torch.max(outputs, 1)
running_loss += loss.item() * inputs.size(0)
n_correct = torch.sum(preds == labels)
running_corrects += n_correct
```

- <https://www.kaggle.com/alexandralorenzo/maskdetection>
 - For YOLO
 - Bounding boxes
 - 588 training images
 - 167 validation images
 - 84 testing images
 - Crowd images

- <https://www.kaggle.com/alexandralorenzo/maskdetection>
 - For YOLO
 - Bounding boxes
 - 588 training images
 - 167 validation images
 - 84 testing images
 - Crowd images

Messy - use as a base!


Datasets

Mask model

- Organize to PyTorch image folder structure
- Collect more data
- Training set
 - 408 images with mask
 - 432 images without mask
- Validation set
 - 80 images with mask
 - 94 images without mask
- Testing set
 - Use own data only!
 - 30 images with mask
 - 29 images without mask
 - Hundreds of tests in practice

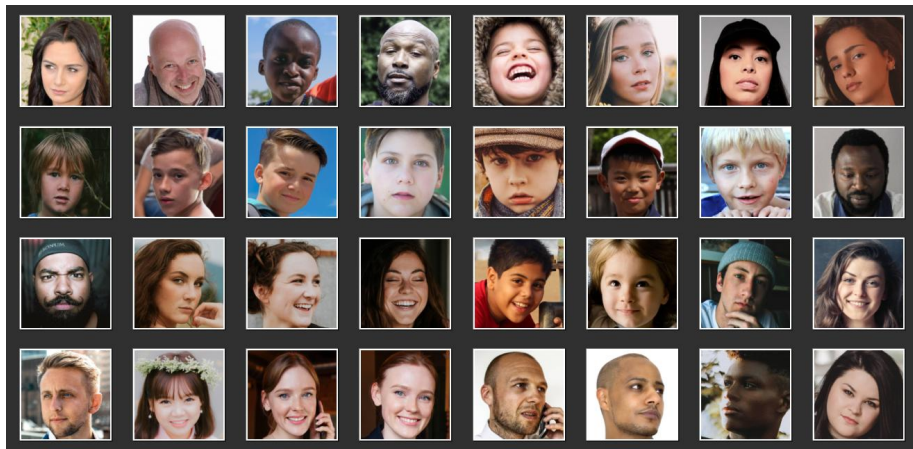
Datasets

Mask model



Datasets

Mask model



- Collect all images with and without masks to another category
- Collect images without faces to another category
- Also custom images without faces
 - Take pictures of floors, inside pockets, walls, ...

- Same training code
- Test on real-world data
- Mask model
 - Around 94% accuracy on test set
 - Very little false positives
- Face model
 - Around 90% accuracy on test set
 - Almost no false positives at all

Model to mobile

Code

- Only a couple lines of code
 - Load models

```
try {  
    face_model = Module.load(assetFilePath(context, "face_model_mobile.pt"));  
    mask_model = Module.load(assetFilePath(context, "mask_model_mobile.pt"));
```

- Predict

```
final Tensor inputTensor = TensorImageUtils.bitmapToFloat32Tensor(image,  
    TensorImageUtils.TORCHVISION_NORM_MEAN_RGB, TensorImageUtils.TORCHVISION_NORM_STD_RGB);  
  
final Tensor outputTensor = mask_model.forward(IValue.from(inputTensor)).toTensor();  
final float[] scores = outputTensor.getDataAsFloatArray();
```

Model to mobile

Code

- Only a couple lines of code

- Load models

```
try {  
    face_model = Module.load(assetFilePath(context, "face_model_mobile.pt"));  
    mask_model = Module.load(assetFilePath(context, "mask_model_mobile.pt"));
```

- Predict

```
final Tensor inputTensor = TensorImageUtils.bitmapToFloat32Tensor(image,  
    TensorImageUtils.TORCHVISION_NORM_MEAN_RGB, TensorImageUtils.TORCHVISION_NORM_STD_RGB);  
final Tensor outputTensor = mask_model.forward(IValue.from(inputTensor)).toTensor();  
final float[] scores = outputTensor.getDataAsFloatArray();
```

- Niclas, our Android expert, will tell you more next!